



High dimensional time series analysis



4. Automatic forecasting algorithms

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models
- 6 Lab Session 9
- 7 Forecast accuracy measures

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models
- 6 Lab Session 9
- 7 Forecast accuracy measures

Historical perspective

- Developed in the 1950s and 1960s as methods (algorithms) to produce point forecasts.
- Combine a “level”, “trend” (slope) and “seasonal” component to describe a time series.
- The rate of change of the components are controlled by “smoothing parameters”: α , β and γ respectively.
- Need to choose best values for the smoothing parameters (and initial states).
- Equivalent ETS state space models developed in the 1990s and 2000s.

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

Perhaps a mix of both?

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

Perhaps a mix of both?

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

How do the level, trend and seasonal components evolve over time?

ETS models

General notation

ETS : ExponenTial Smoothing



Error Trend Season

Error: Additive ("A") or multiplicative ("M")

ETS models

General notation

ETS : ExponenTial Smoothing



Error Trend Season

Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

ETS models

General notation

ETS : ExponenTial Smoothing



Error Trend Season

The diagram shows three arrows pointing upwards from the words 'Error', 'Trend', and 'Season' to the letters 'E', 'T', and 'S' respectively in the 'ETS' part of the notation above.

Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

Seasonality: None ("N"), additive ("A") or multiplicative ("M")

ETS(A,N,N): SES with additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1} + \varepsilon_t$$

State equation

$$\ell_t = \ell_{t-1} + \alpha \varepsilon_t$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

ETS(A,N,N): SES with additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1} + \varepsilon_t$$

State equation

$$\ell_t = \ell_{t-1} + \alpha \varepsilon_t$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

- “innovations” or “single source of error” because equations have the same error process, ε_t .
- Measurement equation: relationship between observations and states.
- Transition/state equation(s): evolution of the state(s) through time.

ETS(M,N,N): SES with multiplicative errors

Forecast equation	$\hat{y}_{T+h T} = \ell_T$
Measurement equation	$y_t = \ell_{t-1}(1 + \varepsilon_t)$
State equation	$\ell_t = \ell_{t-1}(1 + \alpha\varepsilon_t)$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

ETS(M,N,N): SES with multiplicative errors

Forecast equation	$\hat{y}_{T+h T} = \ell_T$
Measurement equation	$y_t = \ell_{t-1}(1 + \varepsilon_t)$
State equation	$\ell_t = \ell_{t-1}(1 + \alpha\varepsilon_t)$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

- Models with additive and multiplicative errors with the same parameters generate the same point forecasts but different prediction intervals.

ETS(A,A,N): Holt's linear trend

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation $y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$

State equations $\ell_t = \ell_{t-1} + b_{t-1} + \alpha\varepsilon_t$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

ETS(A,A,N): Holt's linear trend

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation $y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$

State equations $\ell_t = \ell_{t-1} + b_{t-1} + \alpha\varepsilon_t$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

Multiplicative errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation $y_t = (\ell_{t-1} + b_{t-1})(1 + \varepsilon_t)$

State equations $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha\varepsilon_t)$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

Example: Australian population

```
aus_economy <- global_economy %>% filter(Code == "AUS") %>%  
  mutate(Pop = Population/1e6)  
fit <- aus_economy %>% model(AAN = ETS(Pop))  
report(fit)
```

```
## Series: Pop  
## Model: ETS(A,A,N)  
## Smoothing parameters:  
##   alpha = 1  
##   beta  = 0.327  
##  
## Initial states:  
##   l      b  
## 10.1 0.222  
##  
## sigma^2: 0.0041  
##  
## AIC  AICc  BIC  
## -77.0 -75.8 -66.7
```

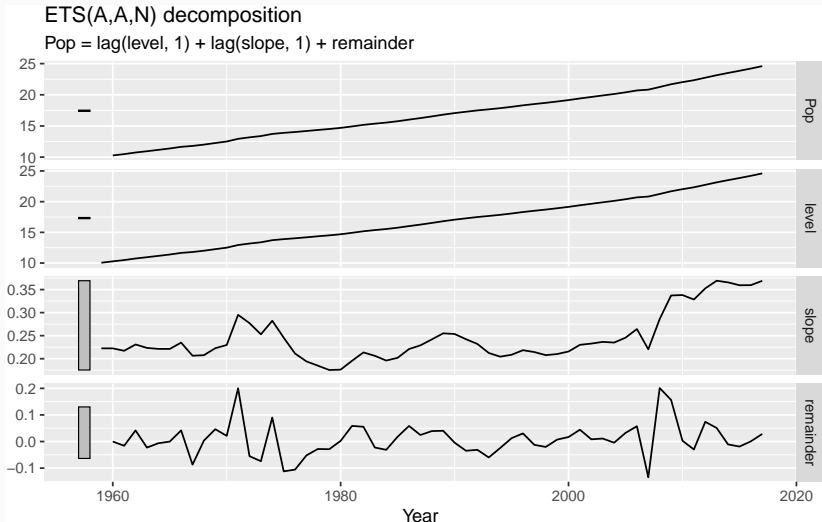
Example: Australian population

```
components(fit)
```

```
## # A dable:                59 x 7 [1Y]
## # Key:                    Country, .model [1]
## # ETS(A,A,N) Decomposition: Pop = lag(level, 1) + lag(slope, 1)
## #   remainder
##   Country   .model Year   Pop level slope remainder
##   <fct>     <chr>  <dbl> <dbl> <dbl> <dbl>      <dbl>
## 1 Australia AAN    1959  NA    10.1 0.222 NA
## 2 Australia AAN    1960  10.3 10.3 0.222 -0.000145
## 3 Australia AAN    1961  10.5 10.5 0.217 -0.0159
## 4 Australia AAN    1962  10.7 10.7 0.231  0.0418
## 5 Australia AAN    1963  11.0 11.0 0.223 -0.0229
## 6 Australia AAN    1964  11.2 11.2 0.221 -0.00641
## 7 Australia AAN    1965  11.4 11.4 0.221 -0.000314
## 8 Australia AAN    1966  11.7 11.7 0.235  0.0418
## 9 Australia AAN    1967  11.8 11.8 0.206 -0.0869
## 10 Australia AAN    1968  12.0 12.0 0.208  0.00350
```

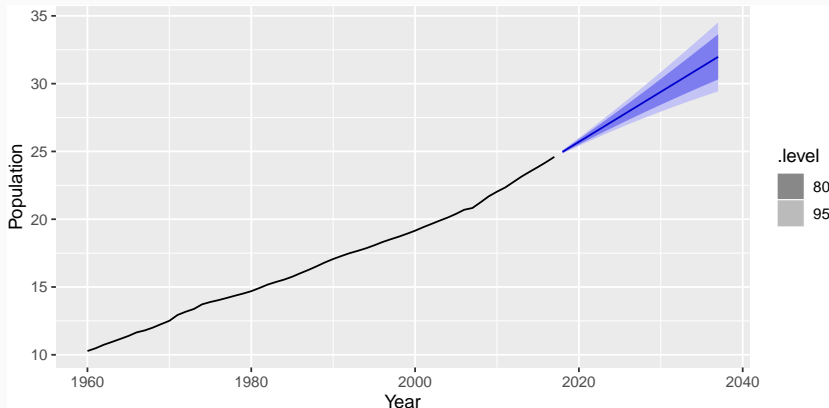
Example: Australian population

```
components(fit) %>% autoplot()
```



Example: Australian population

```
fit %>%  
  forecast(h = 20) %>%  
  autoplot(aus_economy) +  
  ylab("Population") + xlab("Year")
```



ETS(A,Ad,N): Damped trend method

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$

Measurement equation $y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$

State equations $\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$

$$b_t = \phi b_{t-1} + \beta \varepsilon_t$$

ETS(A,Ad,N): Damped trend method

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$

Measurement equation $y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$

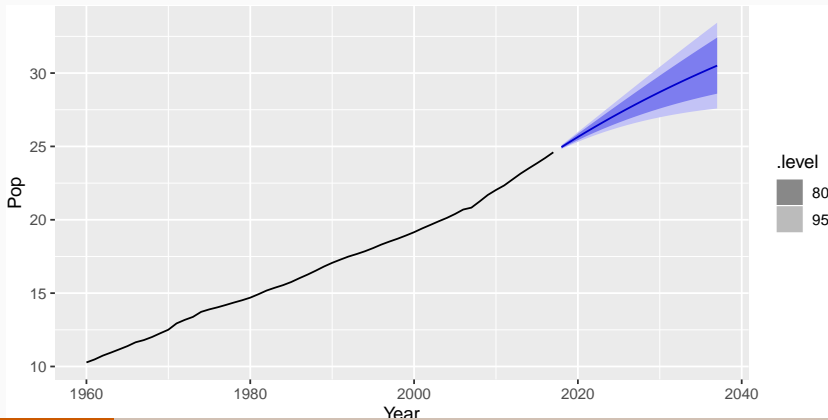
State equations $\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$

$$b_t = \phi b_{t-1} + \beta \varepsilon_t$$

- Damping parameter $0 < \phi < 1$.
- If $\phi = 1$, identical to Holt's linear trend.
- As $h \rightarrow \infty$, $\hat{y}_{T+h|T} \rightarrow \ell_T + \phi b_T / (1 - \phi)$.
- Short-run forecasts trended, long-run forecasts constant.

Example: Australian population

```
aus_economy %>%  
  model(holt = ETS(Pop ~ trend("Ad"))) %>%  
  forecast(h = 20) %>%  
  autoplot(aus_economy)
```



Example: National populations

```
fit <- global_economy %>%  
  mutate(Pop = Population/1e6) %>%  
  model(ets = ETS(Pop))  
fit
```

```
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country      ets  
##   <fct>        <model>  
## 1 Afghanistan <ETS(A,A,N)>  
## 2 Albania     <ETS(M,A,N)>  
## 3 Algeria     <ETS(M,A,N)>  
## 4 American Samoa <ETS(M,A,N)>  
## 5 Andorra     <ETS(M,A,N)>  
## 6 Angola      <ETS(M,A,N)>  
## 7 Antigua and Barbuda <ETS(M,A,N)>  
## 8 Arab World  <ETS(M,A,N)>  
## 9 Argentina   <ETS(A,A,N)>  
## 10 Armenia    <ETS(M,A,N)>  
## # ... with 253 more rows
```

Example: National populations

```
fit %>%  
  forecast(h = 5)
```

```
## # A tibble: 1,315 x 5 [1Y]  
## # Key:      Country, .model [263]  
##   Country      .model Year   Pop .distribution  
##   <fct>        <chr>  <dbl> <dbl> <dist>  
## 1 Afghanistan ets     2018  36.4 N(36, 0.012)  
## 2 Afghanistan ets     2019  37.3 N(37, 0.059)  
## 3 Afghanistan ets     2020  38.2 N(38, 0.164)  
## 4 Afghanistan ets     2021  39.0 N(39, 0.351)  
## 5 Afghanistan ets     2022  39.9 N(40, 0.644)  
## 6 Albania     ets     2018   2.87 N(2.9, 0.00012)  
## 7 Albania     ets     2019   2.87 N(2.9, 0.00060)  
## 8 Albania     ets     2020   2.87 N(2.9, 0.00169)  
## 9 Albania     ets     2021   2.86 N(2.9, 0.00362)
```

ETS(A,A,A): Holt-Winters additive method

Forecast equation $\hat{y}_{t+h|t} = \ell_t + hb_t + s_{t+h-m(k+1)}$

Observation equation $y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$

State equations $\ell_t = \ell_{t-1} + b_{t-1} + \alpha\varepsilon_t$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

$$s_t = s_{t-m} + \gamma\varepsilon_t$$

- $k = \text{integer part of } (h - 1)/m$.
- $\sum_i s_i \approx 0$.
- Parameters: $0 \leq \alpha \leq 1$, $0 \leq \beta \leq 1$, $0 \leq \gamma \leq 1 - \alpha$ and $m = \text{period of seasonality (e.g. } m = 4 \text{ for quarterly data)}$.

ETS(M,A,M): Holt-Winters multiplicative method

Forecast equation $\hat{y}_{t+h|t} = (\ell_t + hb_t)s_{t+h-m(k+1)}$

Observation equation $y_t = (\ell_{t-1} + b_{t-1})s_{t-m}(1 + \varepsilon_t)$

State equations $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha\varepsilon_t)$

$$b_t = b_{t-1}(1 + \beta\varepsilon_t)$$

$$s_t = s_{t-m}(1 + \gamma\varepsilon_t)$$

- k is integer part of $(h - 1)/m$.
- $\sum_i s_i \approx m$.
- Parameters: $0 \leq \alpha \leq 1$, $0 \leq \beta^* \leq 1$, $0 \leq \gamma \leq 1 - \alpha$ and $m =$ period of seasonality (e.g. $m = 4$ for quarterly data).

Example: Australian holiday tourism

```
holidays <- tourism %>%  
  filter(Purpose == "Holiday")  
fit <- holidays %>% model(ets = ETS(Trips))  
fit
```

```
## # A mable: 76 x 4  
## # Key:      Region, State, Purpose [76]  
##   Region                State                Purpose ets  
##   <chr>                 <chr>                 <chr>  <model>  
## 1 Adelaide             South Australia    Holiday <ETS(A,N,A~  
## 2 Adelaide Hills       South Australia    Holiday <ETS(A,A,N~  
## 3 Alice Springs        Northern Territo~ Holiday <ETS(M,N,A~  
## 4 Australia's Coral Coast Western Australia Holiday <ETS(M,N,A~  
## 5 Australia's Golden Outba~ Western Australia Holiday <ETS(M,N,M~  
## 6 Australia's North West Western Australia Holiday <ETS(A,N,A~  
## 7 Australia's South West Western Australia Holiday <ETS(M,N,M~  
## 8 Ballarat             Victoria          Holiday <ETS(M,N,A~  
## 9 Barkly               Northern Territo~ Holiday <ETS(A,N,A~  
## 10 Barossa             South Australia    Holiday <ETS(A,N,N~
```

Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% report()
```

```
## Series: Trips
## Model: ETS(M,N,A)
##   Smoothing parameters:
##     alpha = 0.157
##     gamma = 1e-04
##
##   Initial states:
##     l  s1  s2    s3    s4
## 142 -61 131 -42.2 -27.7
##
##   sigma^2: 0.0388
##
##   AIC AICc  BIC
## 852  854  869
```

Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% components(fit)
```

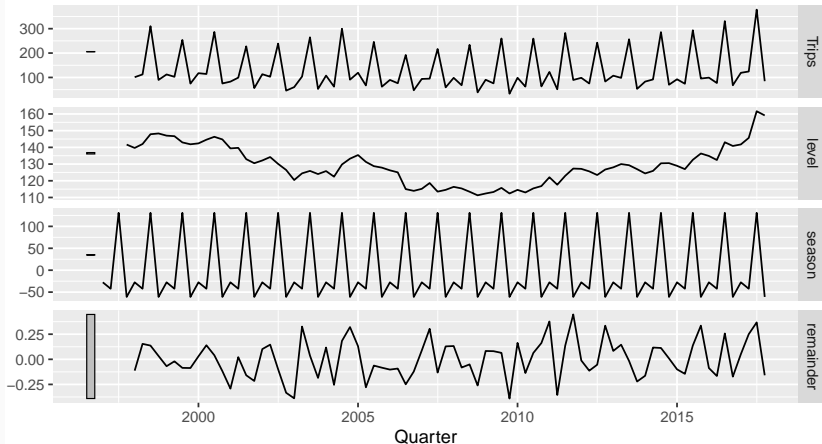
```
## # A dable:                84 x 9 [1Q]
## # Key:                    Region, State, Purpose, .model [1]
## # ETS(M,N,A) Decomposition: Trips = (lag(level, 1) + lag(season,
## #   4)) * (1 + remainder)
##   Region State Purpose .model   Quarter Trips level season
##   <chr>  <chr> <chr>  <chr>      <qtr>  <dbl>  <dbl>  <dbl>
## 1 Snowy~ New ~ Holiday ets      1997 Q1   NA      NA    -27.7
## 2 Snowy~ New ~ Holiday ets      1997 Q2   NA      NA    -42.2
## 3 Snowy~ New ~ Holiday ets      1997 Q3   NA      NA    131.
## 4 Snowy~ New ~ Holiday ets      1997 Q4   NA     142.   -61.0
## 5 Snowy~ New ~ Holiday ets      1998 Q1  101.    140.   -27.7
## 6 Snowy~ New ~ Holiday ets      1998 Q2  112.    142.   -42.2
## 7 Snowy~ New ~ Holiday ets      1998 Q3  310.    148.    131.
## 8 Snowy~ New ~ Holiday ets      1998 Q4   89.8   148.   -61.0
## 9 Snowy~ New ~ Holiday ets      1999 Q1  112.    147.   -27.7
## 10 Snowy~ New ~ Holiday ets      1999 Q2  103.    147.   -42.2
## # ...with 74 more rows, and 1 more variable: remainder <dbl>
```


Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>%  
  components(fit) %>% autoplot()
```

ETS(M,N,A) decomposition

Trips = (lag(level, 1) + lag(season, 4)) * (1 + remainder)



Example: Australian holiday tourism

```
fit %>% forecast()
```

```
## # A tibble: 608 x 7 [1Q]
```

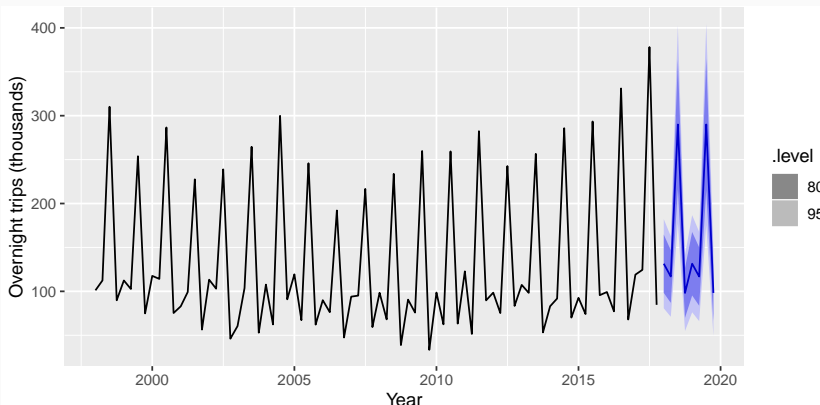
```
## # Key:      Region, State, Purpose, .model [76]
```

##	Region	State	Purpose	.model	Quarter	Trips	.distribution
##	<chr>	<chr>	<chr>	<chr>	<qtr>	<dbl>	<dist>
## 1	Adelaide	South	A~ Holiday	ets	2018 Q1	210.	N(210, 457)
## 2	Adelaide	South	A~ Holiday	ets	2018 Q2	173.	N(173, 473)
## 3	Adelaide	South	A~ Holiday	ets	2018 Q3	169.	N(169, 489)
## 4	Adelaide	South	A~ Holiday	ets	2018 Q4	186.	N(186, 505)
## 5	Adelaide	South	A~ Holiday	ets	2019 Q1	210.	N(210, 521)
## 6	Adelaide	South	A~ Holiday	ets	2019 Q2	173.	N(173, 537)
## 7	Adelaide	South	A~ Holiday	ets	2019 Q3	169.	N(169, 553)
## 8	Adelaide	South	A~ Holiday	ets	2019 Q4	186.	N(186, 569)
## 9	Adelaide~	South	A~ Holiday	ets	2018 Q1	19.4	N(19, 36)
## 10	Adelaide~	South	A~ Holiday	ets	2018 Q2	19.6	N(20, 36)

```
## # ... with 598 more rows
```

Example: Australian holiday tourism

```
fit %>% forecast() %>%  
  filter(Region=="Snowy Mountains") %>%  
  autoplot(holidays) +  
    xlab("Year") + ylab("Overnight trips (thousands)")
```



Exponential smoothing models

Additive Error

Trend Component

Seasonal Component

N (None) A (Additive) M (Multiplicative)

N	(None)	A,N,N	A,N,A	A,N,M
A	(Additive)	A,A,N	A,A,A	A,A,M
A _d	(Additive damped)	A,A _d ,N	A,A _d ,A	A,A_d,M

Multiplicative Error

Trend Component

Seasonal Component

N (None) A (Additive) M (Multiplicative)

N	(None)	M,N,N	M,N,A	M,N,M
A	(Additive)	M,A,N	M,A,A	M,A,M
A _d	(Additive damped)	M,A _d ,N	M,A _d ,A	M,A _d ,M

Estimating ETS models

- Smoothing parameters α , β , γ and ϕ , and the initial states ℓ_0 , b_0 , s_0 , s_{-1} , \dots , s_{-m+1} are estimated by maximising the “likelihood” = the probability of the data arising from the specified model.
- For models with additive errors equivalent to minimising SSE.
- For models with multiplicative errors, **not** equivalent to minimising SSE.

Model selection

Akaike's Information Criterion

$$\text{AIC} = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters initial states estimated in the model.

Model selection

Akaike's Information Criterion

$$\text{AIC} = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters initial states estimated in the model.

Corrected AIC

$$\text{AIC}_c = \text{AIC} + \frac{2(k+1)(k+2)}{T-k}$$

which is the AIC corrected (for small sample bias).

Model selection

Akaike's Information Criterion

$$\text{AIC} = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters initial states estimated in the model.

Corrected AIC

$$\text{AIC}_c = \text{AIC} + \frac{2(k+1)(k+2)}{T-k}$$

which is the AIC corrected (for small sample bias).

Bayesian Information Criterion

$$\text{BIC} = \text{AIC} + k(\log(T) - 2).$$

AIC and cross-validation

Minimizing the AIC assuming Gaussian residuals is asymptotically equivalent to minimizing one-step time series cross validation MSE.

Automatic forecasting

From Hyndman et al. (IJF, 2002):

- 1 Apply each model that is appropriate to the data. Optimize parameters and initial values using MLE.
 - 2 Select best method using AICc.
 - 3 Produce forecasts using best method.
 - 4 Obtain forecast intervals using underlying state space model.
- Method performed very well in M3 competition.
 - Used as a benchmark in the M4 competition.

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models
- 6 Lab Session 9
- 7 Forecast accuracy measures

Find an ETS model for the Gas data from `aus_production`.

- Why is multiplicative seasonality necessary here?
- Experiment with making the trend damped.

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models
- 6 Lab Session 9
- 7 Forecast accuracy measures

ARIMA models

- AR:** autoregressive (lagged observations as inputs)
- I:** integrated (differencing to make series stationary)
- MA:** moving average (lagged errors as inputs)

ARIMA models

AR: autoregressive (lagged observations as inputs)

I: integrated (differencing to make series stationary)

MA: moving average (lagged errors as inputs)

An ARIMA model is rarely interpretable in terms of visible data structures like trend and seasonality. But it can capture a huge range of time series patterns.

Stationarity

Definition

If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

Stationarity

Definition

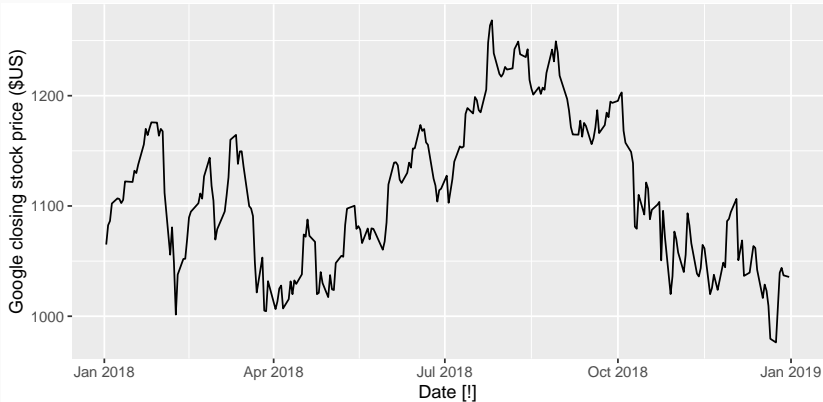
If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

A **stationary series** is:

- roughly horizontal
- constant variance
- no patterns predictable in the long-term

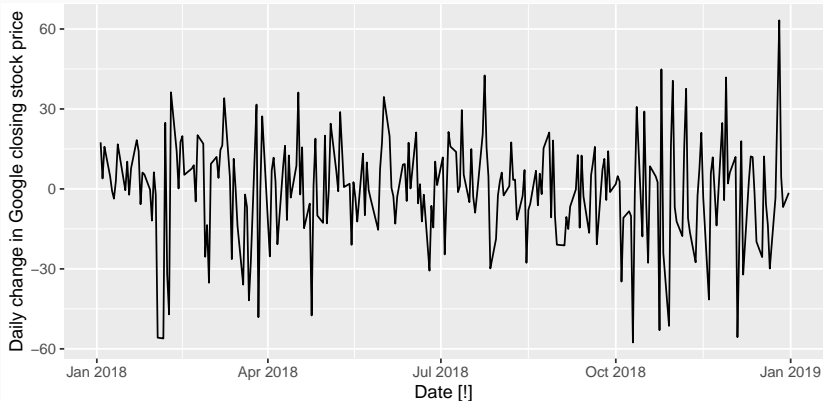
Stationary?

```
gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) == 2018) %>%  
  autoplot(Close) +  
    ylab("Google closing stock price ($US)")
```



Stationary?

```
gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) == 2018) %>%  
  autoplot(difference(Close)) +  
  ylab("Daily change in Google closing stock price")
```



Differencing

- Differencing helps to **stabilize the mean**.
- The differenced series is the *change* between each observation in the original series.
- Occasionally the differenced data will not appear stationary and it may be necessary to difference the data a second time.
- In practice, it is almost never necessary to go beyond second-order differences.

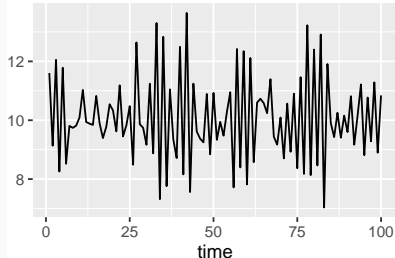
Autoregressive models

Autoregressive (AR) models:

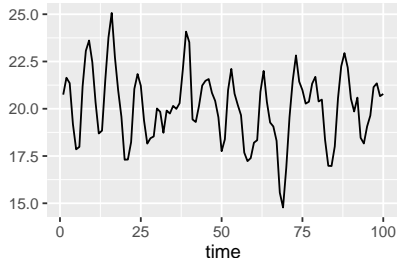
$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$

where ε_t is white noise. This is a multiple regression with **lagged values** of y_t as predictors.

AR(1)



AR(2)



- Cyclic behaviour is possible when $p \geq 2$.

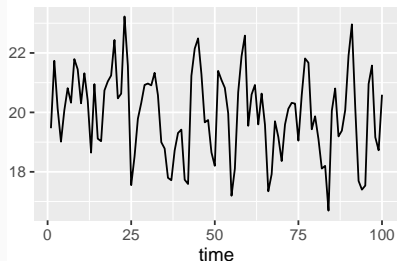
Moving Average (MA) models

Moving Average (MA) models:

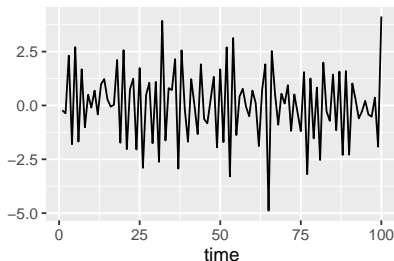
$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$

where ε_t is white noise. This is a multiple regression with **lagged errors** as predictors. *Don't confuse this with moving average smoothing!*

MA(1)



MA(2)



ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of y_t** and **lagged errors**.

ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of y_t** and **lagged errors**.

Autoregressive Integrated Moving Average models

- Combine ARMA model with **differencing**.
- d -differenced series follows an ARMA model.
- Need to choose p, d, q and whether or not to include c .

ARIMA models

ARIMA(p, d, q) model

AR: p = order of the autoregressive part

I: d = degree of first differencing involved

MA: q = order of the moving average part.

- White noise model: ARIMA(0,0,0)
- Random walk: ARIMA(0,1,0) with no constant
- Random walk with drift: ARIMA(0,1,0) with const.
- AR(p): ARIMA($p,0,0$)
- MA(q): ARIMA(0,0, q)

Example: National populations

```
fit <- global_economy %>%  
  model(arima = ARIMA(Population))  
fit
```

```
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country      arima  
##   <fct>        <model>  
## 1 Afghanistan <ARIMA(4,2,1)>  
## 2 Albania     <ARIMA(0,2,2)>  
## 3 Algeria     <ARIMA(2,2,2)>  
## 4 American Samoa <ARIMA(2,2,2)>  
## 5 Andorra     <ARIMA(2,1,2) w/ drift>  
## 6 Angola      <ARIMA(4,2,1)>  
## 7 Antigua and Barbuda <ARIMA(2,1,2) w/ drift>  
## 8 Arab World  <ARIMA(0,2,1)>
```

Example: National populations

```
fit %>% filter(Country=="Australia") %>% report()
```

```
## Series: Population
```

```
## Model: ARIMA(0,2,1)
```

```
##
```

```
## Coefficients:
```

```
##          ma1
```

```
##        -0.661
```

```
## s.e.    0.107
```

```
##
```

```
## sigma^2 estimated as 4.063e+09:  log likelihood=-699
```

```
## AIC=1401   AICc=1402   BIC=1405
```

Example: National populations

```
fit %>% filter(Country=="Australia") %>% report()
```

```
## Series: Population
```

```
## Model: ARIMA(0,2,1)
```

```
##
```

```
## Coefficients:
```

```
##          ma1
```

```
##        -0.661
```

```
## s.e.    0.107
```

```
##
```

```
## sigma^2 estimated as 4.063e+09:  log likelihood=-699
```

```
## AIC=1401   AICc=1402   BIC=1405
```

$$y_t = 2y_{t-1} - y_{t-2} - 0.7\varepsilon_{t-1} + \varepsilon_t$$
$$\varepsilon_t \sim \text{NID}(0, 4 \times 10^9)$$

Understanding ARIMA models

- If $c = 0$ and $d = 0$, the long-term forecasts will go to zero.
- If $c = 0$ and $d = 1$, the long-term forecasts will go to a non-zero constant.
- If $c = 0$ and $d = 2$, the long-term forecasts will follow a straight line.
- If $c \neq 0$ and $d = 0$, the long-term forecasts will go to the mean of the data.
- If $c \neq 0$ and $d = 1$, the long-term forecasts will follow a straight line.
- If $c \neq 0$ and $d = 2$, the long-term forecasts will follow a quadratic trend.

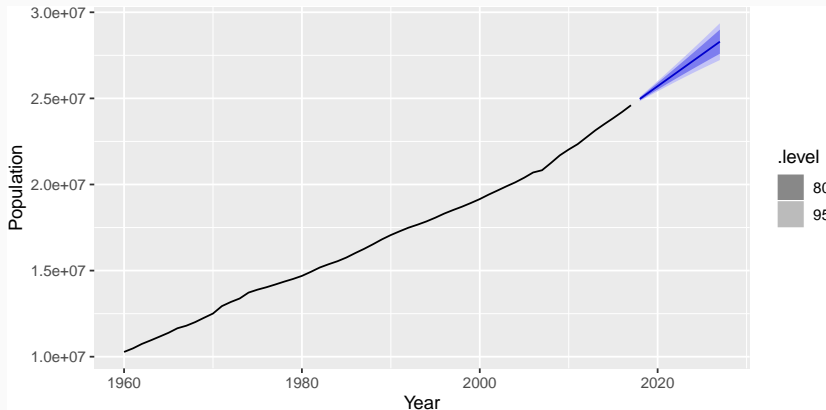
Understanding ARIMA models

Forecast variance and d

- The higher the value of d , the more rapidly the prediction intervals increase in size.
- For $d = 0$, the long-term forecast standard deviation will go to the standard deviation of the historical data.

Example: National populations

```
fit %>% forecast(h=10) %>%  
  filter(Country=="Australia") %>%  
  autoplot(global_economy)
```



How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p , q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

Note: Can't compare AICc for different values of d .

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

How does ARIMA() work?

Step 1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

Step 2: Consider variations of current model:

- vary one of p , q , from current model by ± 1 ;
- p , q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

Step 2: Consider variations of current model:

- vary one of p , q , from current model by ± 1 ;
- p , q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

Repeat Step 2 until no lower AICc can be found.

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models
- 6 Lab Session 9
- 7 Forecast accuracy measures

For the United States GDP data (from `global_economy`):

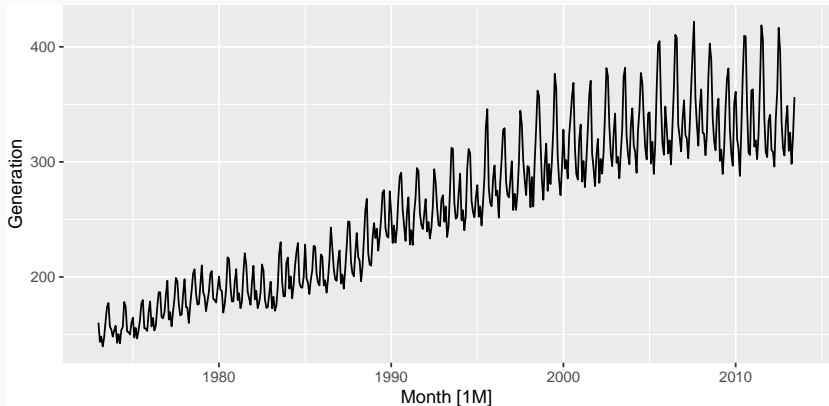
- Fit a suitable ARIMA model for the logged data.
- Produce forecasts of your fitted model. Do the forecasts look reasonable?

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models**
- 6 Lab Session 9
- 7 Forecast accuracy measures

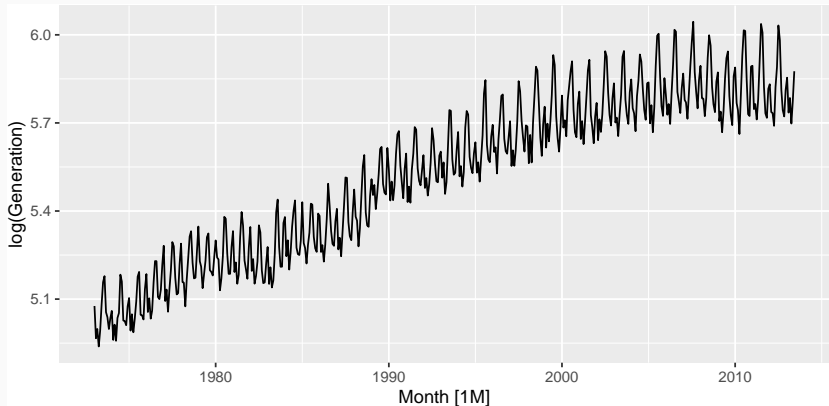
Electricity production

```
usmelec %>% autoplot(  
  Generation  
)
```



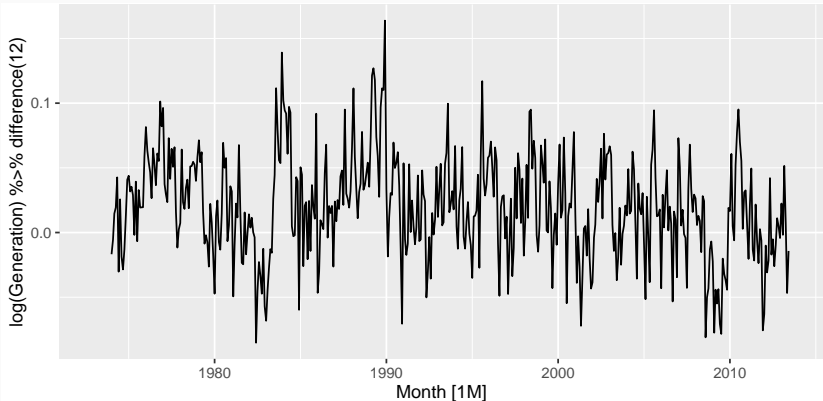
Electricity production

```
usmelec %>% autoplot(  
  log(Generation)  
)
```



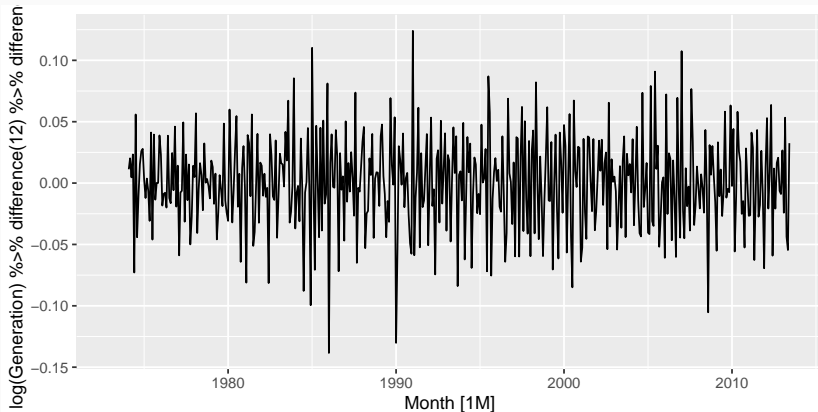
Electricity production

```
usmelec %>% autoplot(  
  log(Generation) %>% difference(12)  
)
```



Electricity production

```
usmelec %>% autoplot(  
  log(Generation) %>% difference(12) %>% difference()  
)
```



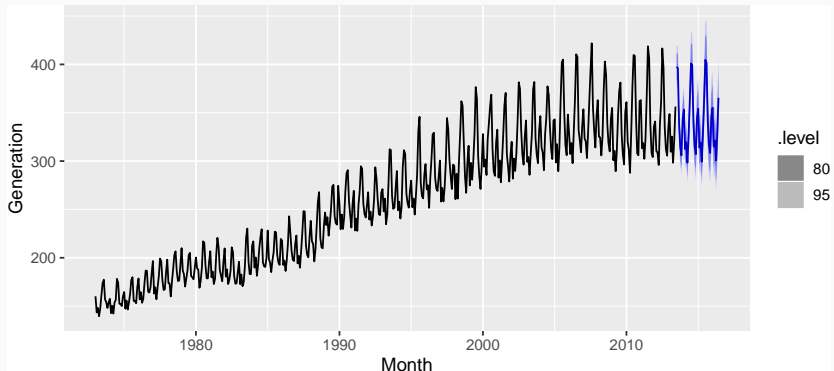
Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  report()
```

```
## Series: Generation  
## Model: ARIMA(1,1,1)(2,1,1)[12]  
## Transformation: log(.x)  
##  
## Coefficients:  
##          ar1          ma1          sar1          sar2          sma1  
##      0.4116   -0.8483    0.0100   -0.1017   -0.8204  
## s.e.  0.0617    0.0348    0.0561    0.0529    0.0357  
##  
## sigma^2 estimated as 0.0006841:  log likelihood=1047  
## AIC=-2082    AICc=-2082    BIC=-2057
```

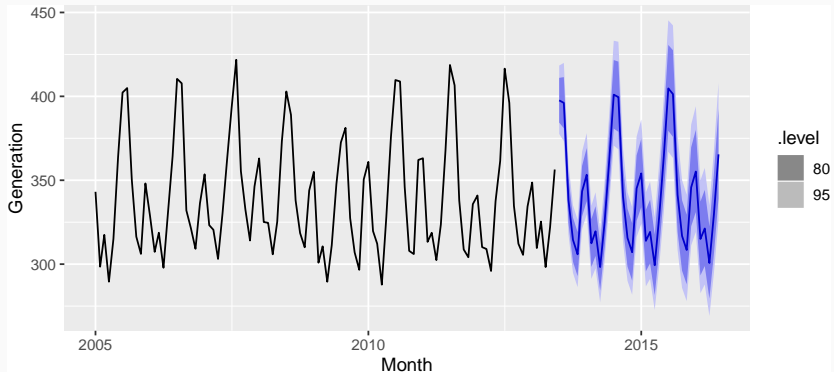
Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  forecast(h="3 years") %>%  
  autoplot(usmelec)
```



Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  forecast(h="3 years") %>%  
  autoplot(filter_index(usmelec, 2005 ~ .))
```



Seasonal ARIMA models

ARIMA	$\underbrace{(p, d, q)}$	$\underbrace{(P, D, Q)_m}$
	↑	↑
	Non-seasonal part of the model	Seasonal part of of the model

- m = number of observations per year.
- d first differences, D seasonal differences
- p AR lags, q MA lags
- P seasonal AR lags, Q seasonal MA lags

Seasonal and non-seasonal terms combine multiplicatively

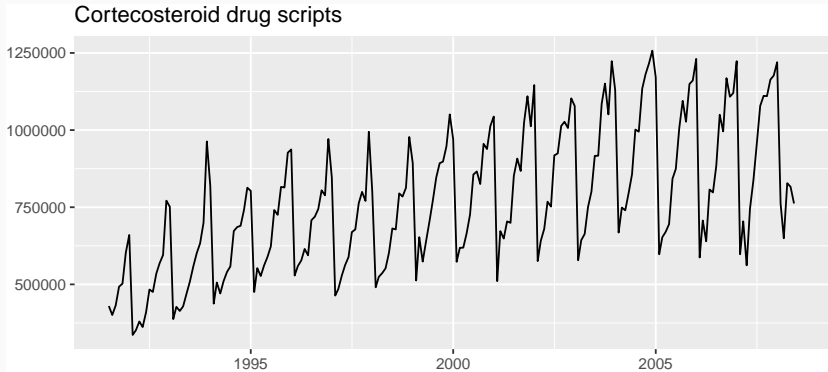
Common ARIMA models

The US Census Bureau uses the following models most often:

ARIMA(0,1,1)(0,1,1) _m	with log transformation
ARIMA(0,1,2)(0,1,1) _m	with log transformation
ARIMA(2,1,0)(0,1,1) _m	with log transformation
ARIMA(0,2,2)(0,1,1) _m	with log transformation
ARIMA(2,1,2)(0,1,1) _m	with no transformation

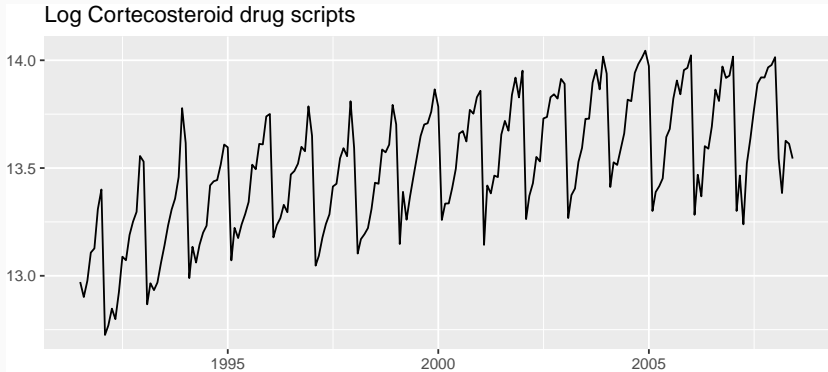
Corticosteroid drug sales

```
h02 <- PBS %>% filter(ATC2 == "H02") %>%  
  summarise(Cost = sum(Cost))  
h02 %>% autoplot(Cost) +  
  xlab("Year") + ylab("") +  
  ggtitle("Corticosteroid drug scripts")
```



Corticosteroid drug sales

```
h02 <- PBS %>% filter(ATC2 == "H02") %>%  
  summarise(Cost = sum(Cost))  
h02 %>% autoplot(log(Cost)) +  
  xlab("Year") + ylab("") +  
  ggtitle("Log Corticosteroid drug scripts")
```



Corticosteroid drug sales

```
fit <- h02 %>%  
  model(auto = ARIMA(log(Cost)))  
report(fit)
```

```
## Series: Cost  
## Model: ARIMA(2,1,0)(0,1,1)[12]  
## Transformation: log(.x)  
##  
## Coefficients:  
##          ar1      ar2      sma1  
##      -0.8491  -0.4207  -0.6401  
## s.e.   0.0712   0.0714   0.0694  
##  
## sigma^2 estimated as 0.004399:  log likelihood=245  
## AIC=-483   AICc=-483   BIC=-470
```

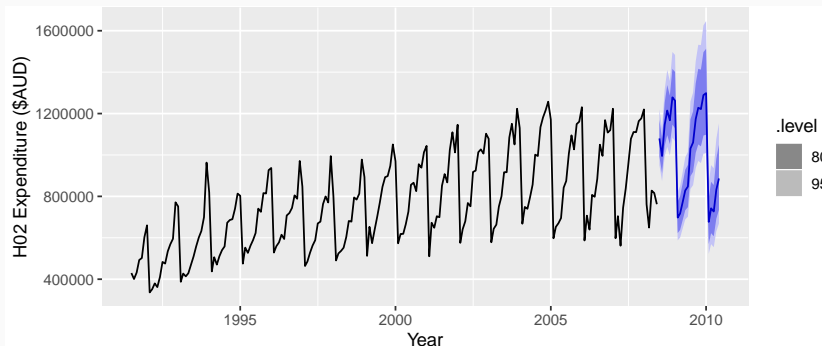
Corticosteroid drug sales

```
fit <- h02 %>%  
  model(best = ARIMA(log(Cost), stepwise = FALSE,  
    approximation = FALSE,  
    order_constraint = p + q + P + Q <= 9))  
report(fit)
```

```
## Series: Cost  
## Model: ARIMA(4,1,1)(2,1,2)[12]  
## Transformation: log(.x)  
##  
## Coefficients:  
##          ar1      ar2      ar3      ar4      ma1      sar1      sar2  
##      -0.0426  0.210   0.202  -0.227  -0.742   0.621  -0.383  
## s.e.   0.2167  0.181   0.114   0.081   0.207   0.242   0.118  
##          sma1      sma2  
##      -1.202   0.496  
## s.e.   0.249   0.214  
##  
## sigma^2 estimated as 0.004061:  log likelihood=254  
## AIC=-489   AICc=-487   BIC=-456
```

Corticosteroid drug sales

```
fit %>% forecast %>% autoplot(h02) +  
  ylab("H02 Expenditure ($AUD)") + xlab("Year")
```



Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models
- 6 Lab Session 9
- 7 Forecast accuracy measures

For the Australian tourism data (from `tourism`):

- Fit a suitable ARIMA model for all data.
- Produce forecasts of your fitted models.
- Check the forecasts for the “Snowy Mountains” and “Melbourne” regions. Do they look reasonable?

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 8
- 5 Seasonal ARIMA models
- 6 Lab Session 9
- 7 Forecast accuracy measures

Training and test sets



- A model which fits the training data well will not necessarily forecast well.
- Forecast accuracy is based only on the test set.

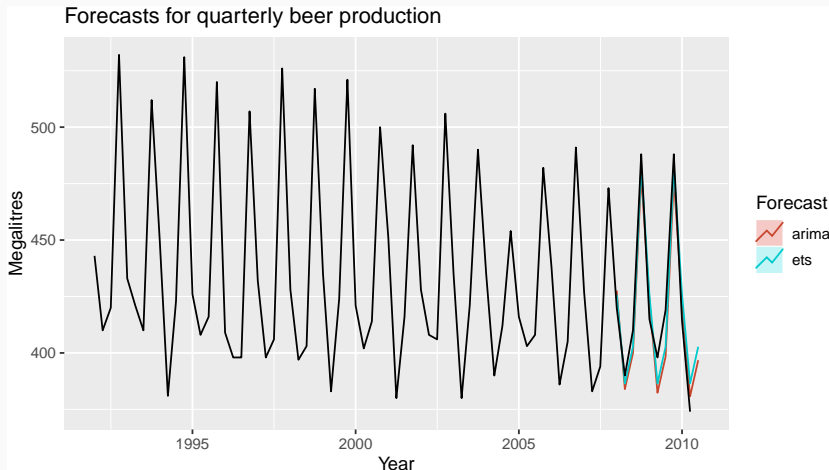
Forecast errors

Forecast “error”: the difference between an observed value and its forecast.

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T},$$

where the training data is given by $\{y_1, \dots, y_T\}$

Measures of forecast accuracy



Measures of forecast accuracy

y_{T+h} = $(T + h)$ th observation, $h = 1, \dots, H$

$\hat{y}_{T+h|T}$ = its forecast based on data up to time T .

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$$

$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

Measures of forecast accuracy

y_{T+h} = $(T + h)$ th observation, $h = 1, \dots, H$

$\hat{y}_{T+h|T}$ = its forecast based on data up to time T .

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$$

$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

- MAE, MSE, RMSE are all scale dependent.
- MAPE is scale independent but is only sensible if $y_t \gg 0$ for all t , and y has a natural zero.

Measures of forecast accuracy

```
recent_production <- aus_production %>%  
  filter(year(Quarter) >= 1992)  
train <- recent_production %>% filter(year(Quarter) <= 2007)  
beer_fit <- train %>%  
  model(  
    ets = ETS(Beer),  
    arima = ARIMA(Beer)  
  )  
beer_fc <- forecast(beer_fit, h="4 years")  
accuracy(beer_fc, aus_production)
```

```
## # A tibble: 2 x 9  
##   .model .type    ME  RMSE   MAE   MPE  MAPE  MASE  ACF1  
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 arima Test   4.18  11.2  10.4  0.940  2.47  0.657  0.145  
## 2 ets   Test   0.854  9.80  8.99  0.151  2.18  0.568  0.207
```

Time series cross-validation

Traditional evaluation

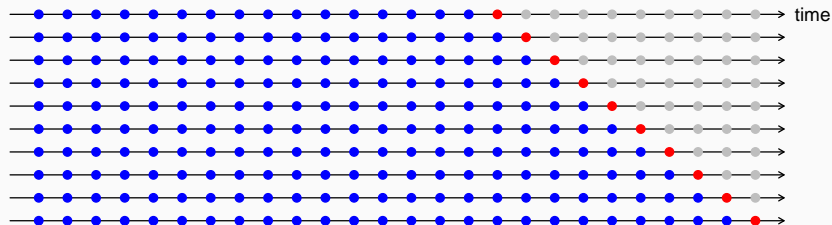


Time series cross-validation

Traditional evaluation



Time series cross-validation

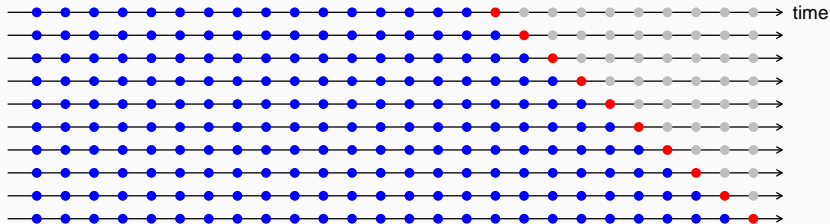


Time series cross-validation

Traditional evaluation



Time series cross-validation



- Forecast accuracy averaged over test sets.
- Also known as “evaluation on a rolling forecasting origin”

Creating the rolling training sets

There are three main rolling types which can be used.

- Stretch: extends a growing length window with new data.
- Slide: shifts a fixed length window through the data.
- Tile: moves a fixed length window without overlap.

Three functions to roll a tsibble: `stretch_tsibble()`, `slide_tsibble()`, and `tile_tsibble()`.

For time series cross-validation, stretching windows are most commonly used.

Creating the rolling training sets

Time series cross-validation

Stretch with a minimum length of 3, growing by 1 each step.

```
beer_stretch <- aus_production %>%  
  stretch_tsibble(.init=1, .step=1)
```

```
## # A tsibble: 23,871 x 8 [1Q]
```

```
## # Key:           .id [218]
```

##	Quarter	Beer	Tobacco	Bricks	Cement	Electricity	Gas	.id
##	<qtr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>
##	1 1956 Q1	284	5225	189	465	3923	5	1
##	2 1956 Q1	284	5225	189	465	3923	5	2
##	3 1956 Q2	213	5178	204	532	4436	6	2
##	4 1956 Q1	284	5225	189	465	3923	5	3
##	5 1956 Q2	213	5178	204	532	4436	6	3
##	6 1956 Q3	227	5297	208	561	4806	7	3
##	7 1956 Q1	284	5225	189	465	3923	5	4
##	8 1956 Q2	213	5178	204	532	4436	6	784
##	9 1956 Q3	227	5297	208	561	4806	7	4

Time series cross-validation

```
fit_cv <- beer_stretch %>%  
  model(arima=ARIMA(Beer), ets=ETS(Beer))
```

```
## # A mable: 218 x 3  
## # Key:      .id [218]  
##      .id arima                                ets  
##      <int> <model>                            <model>  
## 1      1 <NULL model>                          <NULL model>  
## 2      2 <ARIMA(0,1,0) w/ drift>                 <NULL model>  
## 3      3 <ARIMA(0,2,1)>                          <NULL model>  
## 4      4 <ARIMA(0,0,0) w/ mean>                  <ETS(A,N,N)>  
## 5      5 <NULL model>                          <ETS(A,N,N)>  
## 6      6 <NULL model>                          <ETS(A,N,N)>  
## 7      7 <NULL model>                          <ETS(A,N,N)>  
## 8      8 <NULL model>                          <ETS(A,N,N)>  
## 9      9 <ARIMA(0,0,0) (0,1,0) [4]>             <ETS(A,N,N)>  
## 10     10 <ARIMA(0,0,0) (0,1,0) [4]>            <ETS(A,N,N)>
```

Time series cross-validation

Produce one step ahead forecasts from all models.

```
fc_cv <- fit_cv %>% filter(.id>8) %>% forecast(h=1)
```

```
## # A tibble: 420 x 5 [1Q]
## # Key:      .id, .model [420]
##       .id .model Quarter  Beer .distribution
##   <int> <chr>    <qtr> <dbl> <dist>
## 1     9 arima   1958 Q2   228 N(228, 207)
## 2    10 arima   1958 Q3   236 N(236, 177)
## 3    11 arima   1958 Q4   320 N(320, 151)
## 4    12 arima   1959 Q1   272 N(272, 139)
## 5    13 arima   1959 Q2   233 N(233, 137)
## 6    14 arima   1959 Q3   237 N(237, 127)
## 7    15 arima   1959 Q4   313 N(313, 130)
## 8    16 arima   1960 Q1   261 N(261, 120)
## 9    17 arima   1960 Q2   231. N(231, 116)
## 10   18 arima   1960 Q3   242 N(242, 108)
```

Time series cross-validation

```
fc_cv %>% accuracy(aus_production)
```

```
## # A tibble: 2 x 9
```

```
##   .model .type    ME  RMSE   MAE    MPE  MAPE  MASE  
##   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 arima  Test  0.314  16.4  12.9  0.125   3.08  0.834  
## 2 ets    Test  0.199  17.5  13.3  0.0531  3.16  0.855  
0.129
```

A good way to choose the best forecasting model is to find the model with the smallest RMSE computed using time series cross-validation.