



High dimensional time series analysis



4. Automatic forecasting algorithms

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models
- 6 ARIMA vs ETS
- 7 Lab Session 10

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models
- 6 ARIMA vs ETS
- 7 Lab Session 10

Historical perspective

- Developed in the 1950s and 1960s as methods (algorithms) to produce point forecasts.
- Combine a “level”, “trend” (slope) and “seasonal” component to describe a time series.
- The rate of change of the components are controlled by “smoothing parameters”: α , β and γ respectively.
- Need to choose best values for the smoothing parameters (and initial states).
- Equivalent ETS state space models developed in the 1990s and 2000s.

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

Perhaps a mix of both?

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

Perhaps a mix of both?

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

How do the level, trend and seasonal components evolve over time?

ETS models

General notation

ETS : ExponenTial Smoothing



Error Trend Season

Error: Additive ("A") or multiplicative ("M")

ETS models

General notation

ETS : ExponenTial Smoothing



The diagram shows three arrows pointing upwards from the words 'Error', 'Trend', and 'Season' to the letters 'E', 'T', and 'S' respectively in the 'ETS' part of the notation above.

Error Trend Season

Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

ETS models

General notation

ETS : ExponenTial Smoothing



Error Trend Season

The diagram shows three arrows pointing upwards from the words 'Error', 'Trend', and 'Season' to the letters 'E', 'T', and 'S' respectively in the 'ETS' acronym.

Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

Seasonality: None ("N"), additive ("A") or multiplicative ("M")

ETS(A,N,N): SES with additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1} + \varepsilon_t$$

State equation

$$\ell_t = \ell_{t-1} + \alpha \varepsilon_t$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

ETS(A,N,N): SES with additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1} + \varepsilon_t$$

State equation

$$\ell_t = \ell_{t-1} + \alpha \varepsilon_t$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

- “innovations” or “single source of error” because equations have the same error process, ε_t .
- Measurement equation: relationship between observations and states.
- Transition/state equation(s): evolution of the state(s) through time.

ETS(M,N,N): SES with multiplicative errors

Forecast equation	$\hat{y}_{T+h T} = \ell_T$
Measurement equation	$y_t = \ell_{t-1}(1 + \varepsilon_t)$
State equation	$\ell_t = \ell_{t-1}(1 + \alpha\varepsilon_t)$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

ETS(M,N,N): SES with multiplicative errors

Forecast equation	$\hat{y}_{T+h T} = \ell_T$
Measurement equation	$y_t = \ell_{t-1}(1 + \varepsilon_t)$
State equation	$\ell_t = \ell_{t-1}(1 + \alpha\varepsilon_t)$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

- Models with additive and multiplicative errors with the same parameters generate the same point forecasts but different prediction intervals.

ETS(A,A,N): Holt's linear trend

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation $y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$

State equations $\ell_t = \ell_{t-1} + b_{t-1} + \alpha\varepsilon_t$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

ETS(A,A,N): Holt's linear trend

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation $y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$

State equations $\ell_t = \ell_{t-1} + b_{t-1} + \alpha\varepsilon_t$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

Multiplicative errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + hb_T$

Measurement equation $y_t = (\ell_{t-1} + b_{t-1})(1 + \varepsilon_t)$

State equations $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha\varepsilon_t)$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

Example: Australian population

```
aus_economy <- global_economy %>% filter(Code == "AUS") %>%  
  mutate(Pop = Population/1e6)  
fit <- aus_economy %>% model(AAN = ETS(Pop))  
report(fit)
```

```
## Series: Pop  
## Model: ETS(A,A,N)  
## Smoothing parameters:  
##   alpha = 1  
##   beta  = 0.327  
##  
## Initial states:  
##   l      b  
## 10.1 0.222  
##  
## sigma^2: 0.0041  
##  
## AIC AICc BIC  
## -77.0 -75.8 -66.7
```

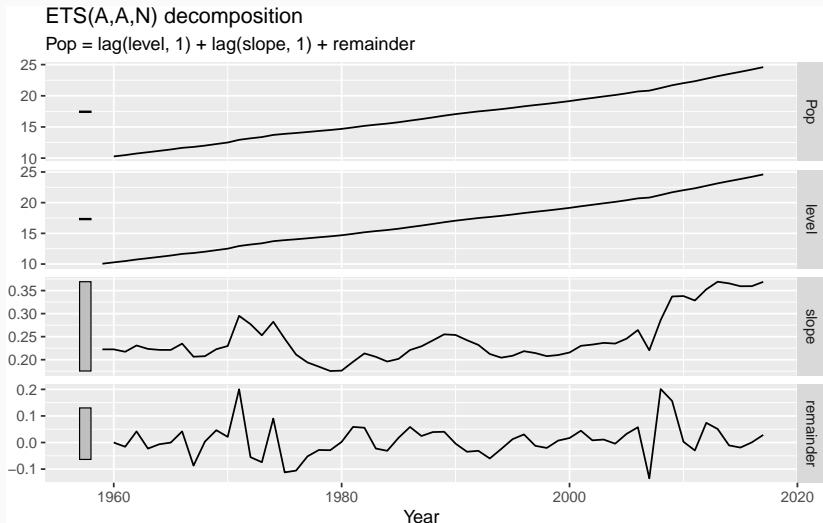
Example: Australian population

```
components(fit)
```

```
## # A dable:                59 x 7 [1Y]
## # Key:                    Country, .model [1]
## # ETS(A,A,N) Decomposition: Pop = lag(level, 1) + lag(slope, 1)
## #   remainder
##   Country   .model Year   Pop level slope remainder
##   <fct>     <chr>  <dbl> <dbl> <dbl> <dbl>      <dbl>
## 1 Australia AAN    1959  NA    10.1 0.222 NA
## 2 Australia AAN    1960  10.3 10.3 0.222 -0.000145
## 3 Australia AAN    1961  10.5 10.5 0.217 -0.0159
## 4 Australia AAN    1962  10.7 10.7 0.231  0.0418
## 5 Australia AAN    1963  11.0 11.0 0.223 -0.0229
## 6 Australia AAN    1964  11.2 11.2 0.221 -0.00641
## 7 Australia AAN    1965  11.4 11.4 0.221 -0.000314
## 8 Australia AAN    1966  11.7 11.7 0.235  0.0418
## 9 Australia AAN    1967  11.8 11.8 0.206 -0.0869
## 10 Australia AAN    1968  12.0 12.0 0.208  0.00350
```

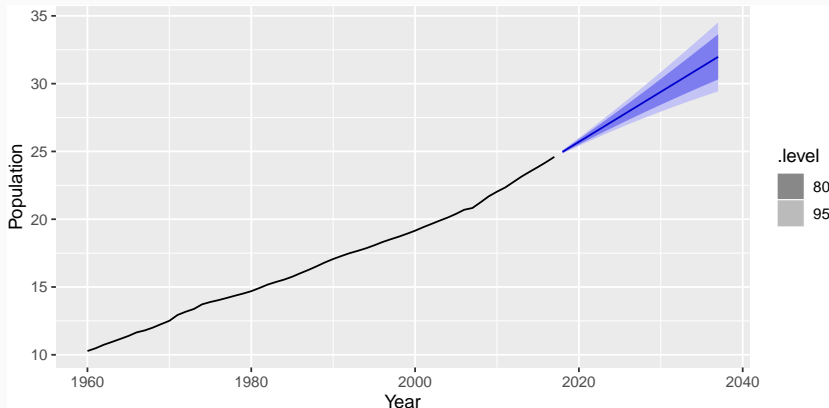
Example: Australian population

```
components(fit) %>% autoplot()
```



Example: Australian population

```
fit %>%  
  forecast(h = 20) %>%  
  autoplot(aus_economy) +  
  ylab("Population") + xlab("Year")
```



ETS(A,Ad,N): Damped trend method

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$

Measurement equation $y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$

State equations $\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$

$$b_t = \phi b_{t-1} + \beta \varepsilon_t$$

ETS(A,Ad,N): Damped trend method

Additive errors

Forecast equation $\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$

Measurement equation $y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$


State equations $\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$

$$b_t = \phi b_{t-1} + \beta \varepsilon_t$$

- Damping parameter $0 < \phi < 1$.
- If $\phi = 1$, identical to Holt's linear trend.
- As $h \rightarrow \infty$, $\hat{y}_{T+h|T} \rightarrow \ell_T + \phi b_T / (1 - \phi)$.
- Short-run forecasts trended, long-run forecasts constant.

Example: Australian population

```
aus_economy %>%  
  model(holt = ETS(Pop ~ trend("Ad"))) %>%  
  forecast(h = 20) %>%  
  autoplot(aus_economy)
```



Example: National populations

```
fit <- global_economy %>%  
  mutate(Pop = Population/1e6) %>%  
  model(ets = ETS(Pop))  
fit
```

```
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country      ets  
##   <fct>        <model>  
## 1 Afghanistan <ETS(A,A,N)>  
## 2 Albania     <ETS(M,A,N)>  
## 3 Algeria     <ETS(M,A,N)>  
## 4 American Samoa <ETS(M,A,N)>  
## 5 Andorra     <ETS(M,A,N)>  
## 6 Angola      <ETS(M,A,N)>  
## 7 Antigua and Barbuda <ETS(M,A,N)>  
## 8 Arab World  <ETS(M,A,N)>  
## 9 Argentina  <ETS(A,A,N)>  
## 10 Armenia    <ETS(M,A,N)>  
## # ... with 253 more rows
```

Example: National populations

```
fit %>%  
  forecast(h = 5)
```

```
## # A tibble: 1,315 x 5 [1Y]  
## # Key:      Country, .model [263]  
##   Country      .model Year   Pop .distribution  
##   <fct>        <chr>   <dbl> <dbl> <dist>  
## 1 Afghanistan ets      2018  36.4 N(36, 0.012)  
## 2 Afghanistan ets      2019  37.3 N(37, 0.059)  
## 3 Afghanistan ets      2020  38.2 N(38, 0.164)  
## 4 Afghanistan ets      2021  39.0 N(39, 0.351)  
## 5 Afghanistan ets      2022  39.9 N(40, 0.644)  
## 6 Albania     ets      2018   2.87 N(2.9, 0.00012)  
## 7 Albania     ets      2019   2.87 N(2.9, 0.00060)  
## 8 Albania     ets      2020   2.87 N(2.9, 0.00169)  
## 9 Albania     ets      2021   2.86 N(2.9, 0.00362)
```

ETS(A,A,A): Holt-Winters additive method

Forecast equation $\hat{y}_{t+h|t} = \ell_t + hb_t + s_{t+h-m(k+1)}$

Observation equation $y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$

State equations $\ell_t = \ell_{t-1} + b_{t-1} + \alpha\varepsilon_t$

$$b_t = b_{t-1} + \beta\varepsilon_t$$

$$s_t = s_{t-m} + \gamma\varepsilon_t$$

- $k = \text{integer part of } (h - 1)/m$.
- $\sum_i s_i \approx 0$.
- Parameters: $0 \leq \alpha \leq 1$, $0 \leq \beta \leq 1$, $0 \leq \gamma \leq 1 - \alpha$ and $m = \text{period of seasonality (e.g. } m = 4 \text{ for quarterly data)}$.

ETS(M,A,M): Holt-Winters multiplicative method

Forecast equation $\hat{y}_{t+h|t} = (\ell_t + hb_t)s_{t+h-m(k+1)}$

Observation equation $y_t = (\ell_{t-1} + b_{t-1})s_{t-m}(1 + \varepsilon_t)$

State equations $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha\varepsilon_t)$

$$b_t = b_{t-1}(1 + \beta\varepsilon_t)$$

$$s_t = s_{t-m}(1 + \gamma\varepsilon_t)$$

- k is integer part of $(h - 1)/m$.
- $\sum_i s_i \approx m$.
- Parameters: $0 \leq \alpha \leq 1$, $0 \leq \beta^* \leq 1$, $0 \leq \gamma \leq 1 - \alpha$ and $m = \text{period of seasonality (e.g. } m = 4 \text{ for quarterly data)}$.

Example: Australian holiday tourism

```
holidays <- tourism %>%  
  filter(Purpose == "Holiday")  
fit <- holidays %>% model(ets = ETS(Trips))  
fit
```

```
## # A mable: 76 x 4
```

```
## # Key:      Region, State, Purpose [76]
```

##	Region	State	Purpose	ets
##	<chr>	<chr>	<chr>	<model>
##	1 Adelaide	South Australia	Holiday	<ETS(A,N,A)>
##	2 Adelaide Hills	South Australia	Holiday	<ETS(A,A,N)>
##	3 Alice Springs	Northern Territory	Holiday	<ETS(M,N,A)>
##	4 Australia's Coral Coast	Western Australia	Holiday	<ETS(M,N,A)>
##	5 Australia's Golden Outback	Western Australia	Holiday	<ETS(M,N,M)>
##	6 Australia's North West	Western Australia	Holiday	<ETS(A,N,A)>
##	7 Australia's South West	Western Australia	Holiday	<ETS(M,N,M)>
##	8 Ballarat	Victoria	Holiday	<ETS(M,N,A)>
##	9 Barkly	Northern Territory	Holiday	<ETS(A,N,A)>
##	10 Barossa	South Australia	Holiday	<ETS(A,N,M)>

Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% report()
```

```
## Series: Trips
## Model: ETS(M,N,A)
## Smoothing parameters:
##   alpha = 0.157
##   gamma = 1e-04
##
## Initial states:
##   l   s1  s2    s3    s4
## 142 -61 131 -42.2 -27.7
##
## sigma^2: 0.0388
##
## AIC AICc BIC
## 852 854 869
```

Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>% components(fit)
```

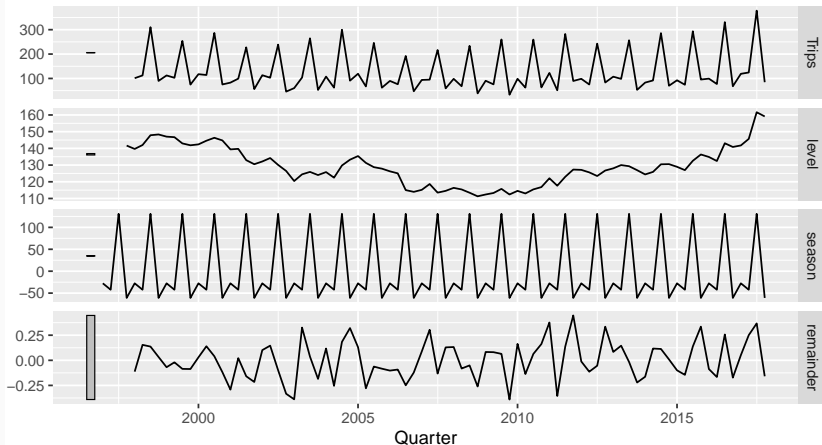
```
## # A dable:                84 x 9 [1Q]
## # Key:                    Region, State, Purpose, .model [1]
## # ETS(M,N,A) Decomposition: Trips = (lag(level, 1) + lag(season, 4))
## # * (1 + remainder)
##   Region State Purpose .model   Quarter Trips level season
##   <chr>  <chr> <chr>  <chr>      <qtr> <dbl> <dbl> <dbl>
## 1 Snowy~ New ~ Holiday ets     1997 Q1  NA      NA    -27.7
## 2 Snowy~ New ~ Holiday ets     1997 Q2  NA      NA    -42.2
## 3 Snowy~ New ~ Holiday ets     1997 Q3  NA      NA    131.
## 4 Snowy~ New ~ Holiday ets     1997 Q4  NA     142.   -61.0
## 5 Snowy~ New ~ Holiday ets     1998 Q1 101.    140.   -27.7
## 6 Snowy~ New ~ Holiday ets     1998 Q2 112.    142.   -42.2
## 7 Snowy~ New ~ Holiday ets     1998 Q3 310.    148.    131.
## 8 Snowy~ New ~ Holiday ets     1998 Q4  89.8   148.   -61.0
## 9 Snowy~ New ~ Holiday ets     1999 Q1 112.    147.   -27.7
## 10 Snowy~ New ~ Holiday ets     1999 Q2 103.    147.   -42.2
## # ...with 74 more rows, and 1 more variable: remainder <dbl>
```


Example: Australian holiday tourism

```
fit %>% filter(Region=="Snowy Mountains") %>%  
  components(fit) %>% autoplot()
```

ETS(M,N,A) decomposition

Trips = (lag(level, 1) + lag(season, 4)) * (1 + remainder)



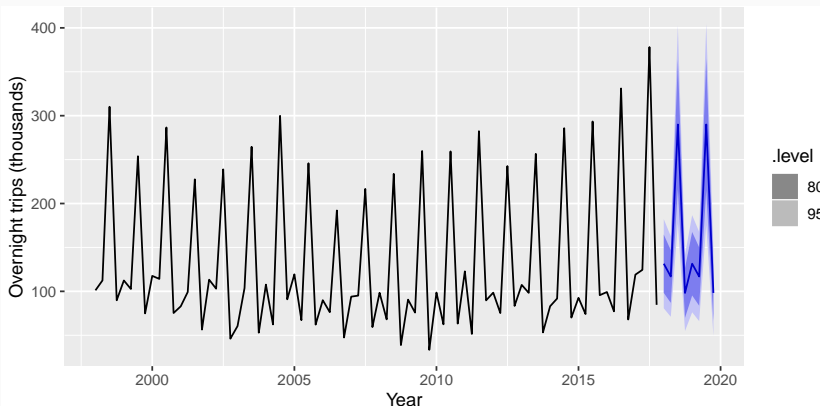
Example: Australian holiday tourism

```
fit %>% forecast()
```

```
## # A tibble: 608 x 7 [1Q]
## # Key:      Region, State, Purpose, .model [76]
##   Region      State      Purpose .model Quarter Trips .distribution
##   <chr>      <chr>      <chr>  <chr>    <qtr> <dbl> <dist>
## 1 Adelaide    South Aus~ Holiday ets    2018 Q1 210. N(210, 457)
## 2 Adelaide    South Aus~ Holiday ets    2018 Q2 173. N(173, 473)
## 3 Adelaide    South Aus~ Holiday ets    2018 Q3 169. N(169, 489)
## 4 Adelaide    South Aus~ Holiday ets    2018 Q4 186. N(186, 505)
## 5 Adelaide    South Aus~ Holiday ets    2019 Q1 210. N(210, 521)
## 6 Adelaide    South Aus~ Holiday ets    2019 Q2 173. N(173, 537)
## 7 Adelaide    South Aus~ Holiday ets    2019 Q3 169. N(169, 553)
## 8 Adelaide    South Aus~ Holiday ets    2019 Q4 186. N(186, 569)
## 9 Adelaide H~ South Aus~ Holiday ets    2018 Q1  19.4 N(19, 36)
## 10 Adelaide H~ South Aus~ Holiday ets    2018 Q2  19.6 N(20, 36)
## # ... with 598 more rows
```

Example: Australian holiday tourism

```
fit %>% forecast() %>%  
  filter(Region=="Snowy Mountains") %>%  
  autoplot(holidays) +  
    xlab("Year") + ylab("Overnight trips (thousands)")
```



Exponential smoothing models

Additive Error

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
Trend Component	N (None)	A,N,N	A,N,A	A,N,M
	A (Additive)	A,A,N	A,A,A	A,A,M
	A _d (Additive damped)	A,A _d ,N	A,A _d ,A	A,A_d,M

Multiplicative Error

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
Trend Component	N (None)	M,N,N	M,N,A	M,N,M
	A (Additive)	M,A,N	M,A,A	M,A,M
	A _d (Additive damped)	M,A _d ,N	M,A _d ,A	M,A _d ,M

Estimating ETS models

- Smoothing parameters α , β , γ and ϕ , and the initial states ℓ_0 , b_0 , s_0 , s_{-1} , \dots , s_{-m+1} are estimated by maximising the “likelihood” = the probability of the data arising from the specified model.
- For models with additive errors equivalent to minimising SSE.
- For models with multiplicative errors, **not** equivalent to minimising SSE.

Model selection

Akaike's Information Criterion

$$\text{AIC} = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters initial states estimated in the model.

Model selection

Akaike's Information Criterion

$$\text{AIC} = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters initial states estimated in the model.

Corrected AIC

$$\text{AIC}_c = \text{AIC} + \frac{2(k+1)(k+2)}{T-k}$$

which is the AIC corrected (for small sample bias).

Model selection

Akaike's Information Criterion

$$\text{AIC} = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters initial states estimated in the model.

Corrected AIC

$$\text{AIC}_c = \text{AIC} + \frac{2(k+1)(k+2)}{T-k}$$

which is the AIC corrected (for small sample bias).

Bayesian Information Criterion

$$\text{BIC} = \text{AIC} + k(\log(T) - 2).$$

AIC and cross-validation

Minimizing the AIC assuming Gaussian residuals is asymptotically equivalent to minimizing one-step time series cross validation MSE.

Automatic forecasting

From Hyndman et al. (IJF, 2002):

- 1 Apply each model that is appropriate to the data. Optimize parameters and initial values using MLE.
 - 2 Select best method using AICc.
 - 3 Produce forecasts using best method.
 - 4 Obtain forecast intervals using underlying state space model.
- Method performed very well in M3 competition.
 - Used as a benchmark in the M4 competition.

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models
- 6 ARIMA vs ETS
- 7 Lab Session 10

Lab Session 7

- 1 Find an ETS model for the Gas data from `aus_production`.
 - ▶ Why is multiplicative seasonality necessary here?
 - ▶ Experiment with making the trend damped.
- 2 Use `ETS()` on some of these series: `tourism`, `gafa_stock`, `pelt`.
 - ▶ Does it always give good forecasts?
 - ▶ Find an example where it does not work well. Can you figure out why?

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models
- 6 ARIMA vs ETS
- 7 Lab Session 10

ARIMA models

- AR:** autoregressive (lagged observations as inputs)
- I:** integrated (differencing to make series stationary)
- MA:** moving average (lagged errors as inputs)

ARIMA models

AR: autoregressive (lagged observations as inputs)

I: integrated (differencing to make series stationary)

MA: moving average (lagged errors as inputs)

An ARIMA model is rarely interpretable in terms of visible data structures like trend and seasonality. But it can capture a huge range of time series patterns.

Stationarity

Definition

If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

Stationarity

Definition

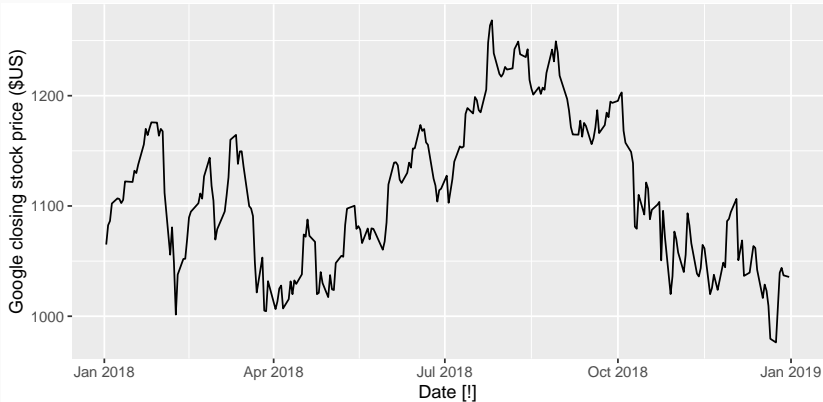
If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

A **stationary series** is:

- roughly horizontal
- constant variance
- no patterns predictable in the long-term

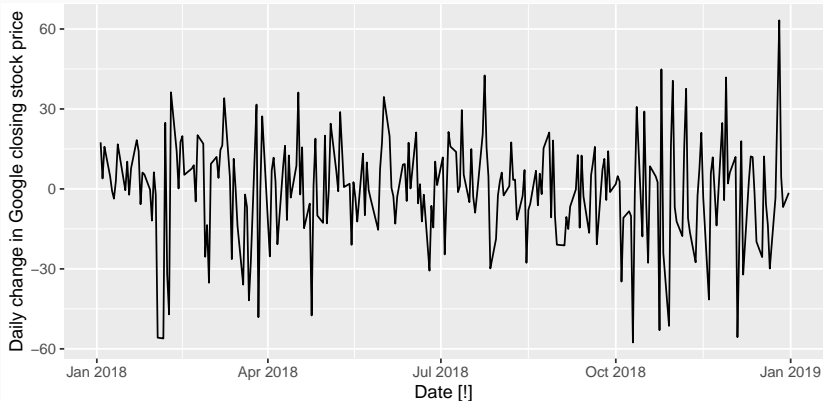
Stationary?

```
gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) == 2018) %>%  
  autoplot(Close) +  
    ylab("Google closing stock price ($US)")
```



Stationary?

```
gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) == 2018) %>%  
  autoplot(difference(Close)) +  
  ylab("Daily change in Google closing stock price")
```



Differencing

- Differencing helps to **stabilize the mean**.
- The differenced series is the *change* between each observation in the original series.
- Occasionally the differenced data will not appear stationary and it may be necessary to difference the data a second time.
- In practice, it is almost never necessary to go beyond second-order differences.

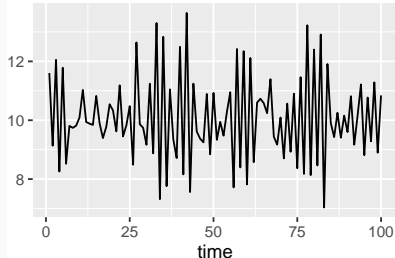
Autoregressive models

Autoregressive (AR) models:

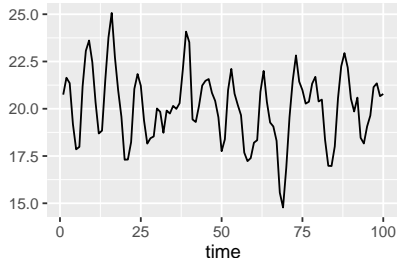
$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$

where ε_t is white noise. This is a multiple regression with **lagged values** of y_t as predictors.

AR(1)



AR(2)



- Cyclic behaviour is possible when $p \geq 2$.

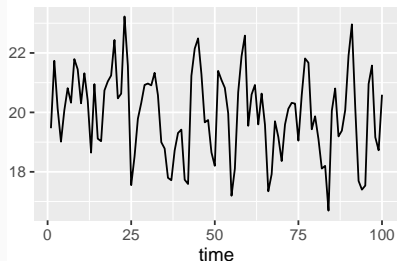
Moving Average (MA) models

Moving Average (MA) models:

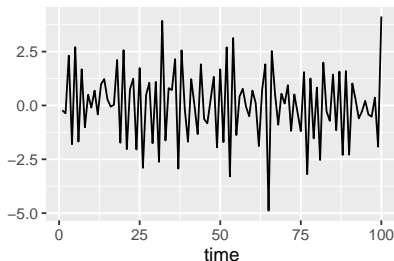
$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$

where ε_t is white noise. This is a multiple regression with **lagged errors** as predictors. *Don't confuse this with moving average smoothing!*

MA(1)



MA(2)



ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of y_t** and **lagged errors**.

ARIMA models

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of y_t** and **lagged errors**.

Autoregressive Integrated Moving Average models

- Combine ARMA model with **differencing**.
- d -differenced series follows an ARMA model.
- Need to choose p, d, q and whether or not to include c .

ARIMA models

ARIMA(p, d, q) model

AR: p = order of the autoregressive part


I: d = degree of first differencing involved

MA: q = order of the moving average part.

- White noise model: ARIMA(0,0,0)
- Random walk: ARIMA(0,1,0) with no constant
- Random walk with drift: ARIMA(0,1,0) with const.
- AR(p): ARIMA($p,0,0$)
- MA(q): ARIMA(0,0, q)


Annual lynx trappings

```
pelt %>% autoplot(Lynx) +  
  xlab("Year") + ylab("Number trapped") +  
  ggtitle("Annual Canadian Lynx Trappings")
```



Annual lynx trappings

```
pelt %>% autoplot(sqrt(Lynx)) +  
  xlab("Year") + ylab("Number trapped") +  
  ggtitle("Annual Canadian Lynx Trappings")
```



Annual lynx trappings

```
pelt %>%  
  model(lynx = ARIMA(sqrt(Lynx))) %>%  
  report()
```

```
## Series: Lynx  
## Model: ARIMA(2,0,1) w/ mean  
## Transformation: sqrt(.x)  
##  
## Coefficients:  
##           ar1          ar2          ma1    constant  
##          1.5059   -0.8645   -0.331         56.33  
## s.e.    0.0583    0.0522    0.108         1.56  
##  
## sigma^2 estimated as 507.9:  log likelihood=-412  
## AIC=834    AICc=835    BIC=847
```

Annual lynx trappings

```
pelt %>%  
  model(lynx = ARIMA(sqrt(Lynx))) %>%  
  report()
```

```
## Series: Lynx
```

```
## Model: ARIMA(2,0,1) w/ mean
```

```
## Transformation: sqrt(
```

$$y_t = \sqrt{(\text{lynx in year } t)}$$

```
##
```

$$y_t = 56.3 + 1.5y_{t-1} - 0.9y_{t-2} - 0.3\varepsilon_{t-1} + \varepsilon_t$$

```
## Coefficients:
```

$$\varepsilon_t \sim \text{NID}(0, 508)$$

```
##           ar1           ar2           ma1    constant
```

```
##           1.5059    -0.8645    -0.331           56.33
```

```
## s.e.    0.0583     0.0522     0.108           1.56
```


```
##
```

```
## sigma^2 estimated as 507.9:  log likelihood=-412
```

```
## AIC=834    AICc=835    BIC=847
```

Annual lynx trappings

```
pelt %>%  
  model(lynx = ARIMA(sqrt(Lynx))) %>%  
  forecast(h=20) %>%  
  autoplot(pelt)
```



Example: National populations

```
fit <- global_economy %>%  
  model(arima = ARIMA(Population))  
fit
```

```
## # A mable: 263 x 2  
## # Key:      Country [263]  
##   Country      arima  
##   <fct>        <model>  
## 1 Afghanistan <ARIMA(4,2,1)>  
## 2 Albania     <ARIMA(0,2,2)>  
## 3 Algeria     <ARIMA(2,2,2)>  
## 4 American Samoa <ARIMA(2,2,2)>  
## 5 Andorra     <ARIMA(2,1,2) w/ drift>  
## 6 Angola      <ARIMA(4,2,1)>  
## 7 Antigua and Barbuda <ARIMA(2,1,2) w/ drift>  
## 8 Arab World  <ARIMA(0,2,1)>
```


Example: National populations

```
fit %>% filter(Country=="Australia") %>% report()
```

```
## Series: Population
## Model: ARIMA(0,2,1)
##
## Coefficients:
##          ma1
##        -0.661
## s.e.    0.107
##
## sigma^2 estimated as 4.063e+09:  log likelihood=-699
## AIC=1401   AICc=1402   BIC=1405
```

Example: National populations

```
fit %>% filter(Country=="Australia") %>% report()
```

```
## Series: Population
```

```
## Model: ARIMA(0,2,1)
```

```
##
```

```
## Coefficients:
```

```
##          ma1
```

```
##        -0.661
```

```
## s.e.    0.107
```

```
##
```

```
## sigma^2 estimated as 4.063e+09:  log likelihood=-699
```

```
## AIC=1401   AICc=1402   BIC=1405
```

$$y_t = 2y_{t-1} - y_{t-2} - 0.7\varepsilon_{t-1} + \varepsilon_t$$
$$\varepsilon_t \sim \text{NID}(0, 4 \times 10^9)$$

Understanding ARIMA models

- If $c = 0$ and $d = 0$, the long-term forecasts will go to zero.
- If $c = 0$ and $d = 1$, the long-term forecasts will go to a non-zero constant.
- If $c = 0$ and $d = 2$, the long-term forecasts will follow a straight line.
- If $c \neq 0$ and $d = 0$, the long-term forecasts will go to the mean of the data.
- If $c \neq 0$ and $d = 1$, the long-term forecasts will follow a straight line.
- If $c \neq 0$ and $d = 2$, the long-term forecasts will follow a quadratic trend.


Understanding ARIMA models

Forecast variance and d

- The higher the value of d , the more rapidly the prediction intervals increase in size.
- For $d = 0$, the long-term forecast standard deviation will go to the standard deviation of the historical data.

Example: National populations

```
fit %>% forecast(h=10) %>%  
  filter(Country=="Australia") %>%  
  autoplot(global_economy)
```



How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p , q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p+q+k+1) \left[1 + \frac{(p+q+k+2)}{T-p-q-k-2} \right].$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

Note: Can't compare AICc for different values of d .

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

Step 2: Consider variations of current model:

- vary one of p , q , from current model by ± 1 ;
- p , q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d , 2)

ARIMA(0, d , 0)

ARIMA(1, d , 0)

ARIMA(0, d , 1)

Step 2: Consider variations of current model:

- vary one of p , q , from current model by ± 1 ;
- p , q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

Repeat Step 2 until no lower AICc can be found.

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models
- 6 ARIMA vs ETS
- 7 Lab Session 10

For the United States GDP data (from `global_economy`):


- Fit a suitable ARIMA model for the logged data.
- Produce forecasts of your fitted model. Do the forecasts look reasonable?

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models**
- 6 ARIMA vs ETS
- 7 Lab Session 10


Electricity production

```
usmelec %>% autoplot(  
  Generation  
)
```




Electricity production

```
usmelec %>% autoplot(  
  log(Generation)  
)
```




Electricity production

```
usmelec %>% autoplot(  
  log(Generation) %>% difference(12)  
)
```




Electricity production

```
usmelec %>% autoplot(  
  log(Generation) %>% difference(12) %>% difference()  
)
```




Example: US electricity production

```
usmelec %>% autoplot(Generation)
```



Example: US electricity production

```
usmelec %>% autoplot(log(Generation))
```




Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  report()
```

```
## Series: Generation  
## Model: ARIMA(1,1,1)(2,1,1)[12]  
## Transformation: log(.x)  
##  
## Coefficients:  
##          ar1          ma1          sar1          sar2          sma1  
##      0.4116   -0.8483    0.0100   -0.1017   -0.8204  
## s.e.  0.0617    0.0348    0.0561    0.0529    0.0357  
##  
## sigma^2 estimated as 0.0006841:  log likelihood=1047  
## AIC=-2082    AICc=-2082    BIC=-2057
```


Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  forecast(h="3 years") %>%  
  autoplot(usmelec)
```



Example: US electricity production

```
usmelec %>%  
  model(arima = ARIMA(log(Generation))) %>%  
  forecast(h="3 years") %>%  
  autoplot(filter_index(usmelec, 2005 ~ .))
```



Seasonal ARIMA models

ARIMA	$\underbrace{(p, d, q)}$	$\underbrace{(P, D, Q)_m}$
	↑	↑
	Non-seasonal part of the model	Seasonal part of of the model

where m = number of observations per year.

- d first differences
- D seasonal differences
- p AR lags
- q MA lags


Common ARIMA models

The US Census Bureau uses the following models most often:

ARIMA(0,1,1)(0,1,1) _m	with log transformation
ARIMA(0,1,2)(0,1,1) _m	with log transformation
ARIMA(2,1,0)(0,1,1) _m	with log transformation
ARIMA(0,2,2)(0,1,1) _m	with log transformation
ARIMA(2,1,2)(0,1,1) _m	with no transformation


Corticosteroid drug sales

```
h02 %>%  
  mutate(log(Cost)) %>%  
  gather() %>%  
  ggplot(aes(x = Month, y = value)) +  
  geom_line() +  
  facet_grid(key ~ ., scales = "free_y") +  
  xlab("Year") + ylab("") +  
  ggtitle("Corticosteroid drug scripts (H02)")
```



Corticosteroid drug sales

```
h02 %>% gg_tsdisplay(difference(log(Cost),12),
```



Corticosteroid drug sales

- Choose $D = 1$ and $d = 0$.
- Spikes in PACF at lags 12 and 24 suggest seasonal AR(2) term.
- Spikes in PACF suggests possible non-seasonal AR(3) term.
- Initial candidate model: $\text{ARIMA}(3,0,0)(2,1,0)_{12}$.

Corticosteroid drug sales

.model	AICc
ARIMA(3,0,1)(0,1,2)[12]	-485
ARIMA(3,0,1)(1,1,1)[12]	-484
ARIMA(3,0,1)(0,1,1)[12]	-484
ARIMA(3,0,1)(2,1,0)[12]	-476
ARIMA(3,0,0)(2,1,0)[12]	-475
ARIMA(3,0,2)(2,1,0)[12]	-475
ARIMA(3,0,1)(1,1,0)[12]	-463

Corticosteroid drug sales

```
fit <- h02 %>%  
  model(best = ARIMA(log(Cost) ~ 0 + pdq(3,0,1) + PDQ(0,1,2)))  
report(fit)
```

```
## Series: Cost
```

```
## Model: ARIMA(3,0,1)(0,1,2)[12]
```

```
## Transformation: log(.x)
```

```
##
```

```
## Coefficients:
```

```
##          ar1          ar2          ar3          ma1          sma1          sma2
```

```
##        -0.160    0.5481    0.5678    0.383    -0.5222    -0.1769
```

```
## s.e.      0.164    0.0878    0.0942    0.190     0.0861     0.0872
```


```
##
```

```
## sigma^2 estimated as 0.004289: log likelihood=250
```

```
## AIC=-486    AICc=-485    BIC=-463
```

Corticosteroid drug sales

```
augment(fit) %>%  
  gg_tsdisplay(.resid, lag_max=36, plot_type = "hist")
```



Corticosteroid drug sales

```
augment(fit) %>%  
  features(.resid, ljung_box, lag = 36, dof = 6)
```

```
## # A tibble: 1 x 3  
##   .model lb_stat lb_pvalue  
##   <chr>    <dbl>    <dbl>  
## 1 best      50.5      0.0109
```

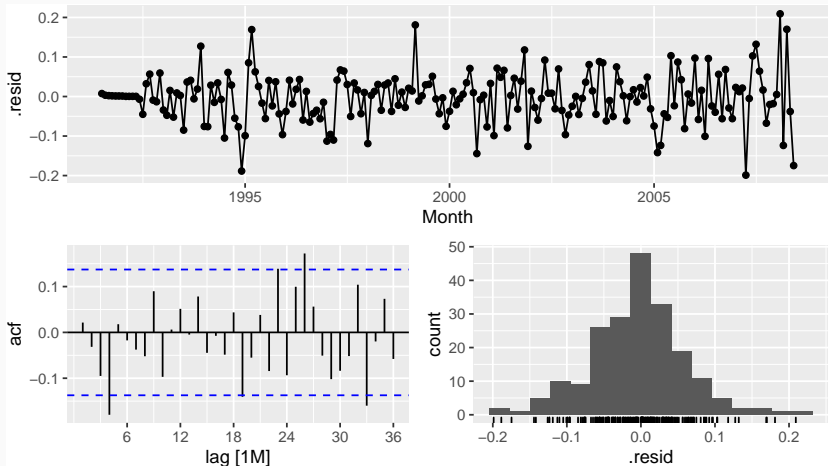

Corticosteroid drug sales

```
fit <- h02 %>% model(auto = ARIMA(log(Cost)))  
report(fit)
```

```
## Series: Cost  
## Model: ARIMA(2,1,0)(0,1,1)[12]  
## Transformation: log(.x)  
##  
## Coefficients:  
##           ar1      ar2      sma1  
##      -0.8491  -0.4207  -0.6401  
## s.e.   0.0712   0.0714   0.0694  
##  
## sigma^2 estimated as 0.004399:  log likelihood=245  
## AIC=-483   AICc=-483   BIC=-470
```

Corticosteroid drug sales

```
augment(fit) %>%  
  gg_tsdisplay(.resid, lag_max = 36, plot_type = "hist")
```



Corticosteroid drug sales

```
augment(fit) %>%  
  features(.resid, ljung_box, lag = 36, dof = 5)
```

```
## # A tibble: 1 x 3  
##   .model lb_stat lb_pvalue  
##   <chr>    <dbl>    <dbl>  
## 1 auto      57.5    0.00260
```


Corticosteroid drug sales

```
fit <- h02 %>%  
  model(best = ARIMA(log(Cost), stepwise = FALSE,  
    approximation = FALSE,  
    order_constraint = p + q + P + Q <= 9))  
report(fit)
```

```
## Series: Cost  
## Model: ARIMA(4,1,1)(2,1,2)[12]  
## Transformation: log(.x)  
##  
## Coefficients:  
##          ar1      ar2      ar3      ar4      ma1      sar1      sar2      sma1  
##      -0.0426  0.210   0.202  -0.227  -0.742   0.621  -0.383   -  
1.202  
## s.e.    0.2167  0.181   0.114   0.081   0.207   0.242   0.118   0.249  
##          sma2  
##        0.496  
## s.e.    0.214  
##  
## sigma^2 estimated as 0.004061:  log likelihood=254  
## AIC=-489   AICc=-487   BIC=-456
```

Corticosteroid drug sales

```
augment(fit) %>%  
  gg_tsdisplay(.resid, lag_max = 36, plot_type = "hist")
```



Corticosteroid drug sales

```
augment(fit) %>%  
  features(.resid, ljung_box, lag = 36, dof = 9)
```

```
## # A tibble: 1 x 3  
##   .model lb_stat lb_pvalue  
##   <chr>    <dbl>    <dbl>  
## 1 best      35.1      0.136
```

Corticosteroid drug sales

Training data: July 1991 to June 2006

Test data: July 2006–June 2008

```
fit <- h02 %>%  
  filter_index(~ "2006 Jun") %>%  
  model(  
    ARIMA(log(Cost) ~ pdq(3, 0, 0) + PDQ(2, 1, 0)),  
    ARIMA(log(Cost) ~ pdq(3, 0, 1) + PDQ(2, 1, 0)),  
    ARIMA(log(Cost) ~ pdq(3, 0, 2) + PDQ(2, 1, 0)),  
    ARIMA(log(Cost) ~ pdq(3, 0, 1) + PDQ(1, 1, 0))  
    # ... #  
  )  
  
fit %>%  
  forecast(h = "2 years") %>%  
  accuracy(h02 %>% filter_index("2006 Jul" ~ .))
```

Corticosteroid drug sales


```
models <- list(  
  c(3,0,0,2,1,0),  
  c(3,0,1,2,1,0),  
  c(3,0,2,2,1,0),  
  c(3,0,1,1,1,0),  
  c(3,0,1,0,1,1),  
  c(3,0,1,0,1,2),  
  c(3,0,1,1,1,1),  
  c(3,0,3,0,1,1),  
  c(3,0,2,0,1,1),  
  c(2,1,3,0,1,1),  
  c(2,1,4,0,1,1),  
  c(2,1,5,0,1,1),  
  c(4,1,1,2,1,2))
```


Corticosteroid drug sales

- Models with lowest AICc values tend to give slightly better results than the other models.
- AICc comparisons must have the same orders of differencing. But RMSE test set comparisons can involve any models.
- Use the best model available, even if it does not pass all tests.

Corticosteroid drug sales

```
fit <- h02 %>%  
  model(ARIMA(Cost ~ 0 + pdq(3,0,1) + PDQ(0,1,2)))  
fit %>% forecast %>% autoplot(h02) +  
  ylab("H02 Expenditure ($AUD)") + xlab("Year")
```



Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models
- 6 ARIMA vs ETS
- 7 Lab Session 10

ARIMA vs ETS

- Myth that ARIMA models are more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.
- Many ARIMA models have no exponential smoothing counterparts.
- ETS models all non-stationary. Models with seasonality or non-damped trend (or both) have two unit roots; all other models have one unit root.

Equivalences

ETS model	ARIMA model	Parameters
ETS(A,N,N)	ARIMA(0,1,1)	$\theta_1 = \alpha - 1$
ETS(A,A,N)	ARIMA(0,2,2)	$\theta_1 = \alpha + \beta - 2$ $\theta_2 = 1 - \alpha$
ETS(A,A,N)	ARIMA(1,1,2)	$\phi_1 = \phi$ $\theta_1 = \alpha + \phi\beta - 1 - \phi$ $\theta_2 = (1 - \alpha)\phi$
ETS(A,N,A)	ARIMA(0,0,m)(0,1,0) _m	
ETS(A,A,A)	ARIMA(0,1,m + 1)(0,1,0) _m	
ETS(A,A,A)	ARIMA(1,0,m + 1)(0,1,0) _m	

Outline

- 1 Exponential smoothing
- 2 Lab Session 7
- 3 ARIMA models
- 4 Lab Session 9
- 5 Seasonal ARIMA models
- 6 ARIMA vs ETS
- 7 Lab Session 10

Lab Session 10

For the `fma::condmilk` series:

- Do the data need transforming? If so, find a suitable transformation.
- Are the data stationary? If not, find an appropriate differencing which yields stationary data.
- Identify a couple of ARIMA models that might be useful in describing the time series.
- Which of your models is the best according to their AIC values?
- Estimate the parameters of your best model and do diagnostic testing on the residuals. Do the residuals resemble white noise? If not, try to find another ARIMA model which fits better.
- Forecast the next 24 months of data using your preferred model.